

Basic Approaches for Missing Values

Basic Approaches for Missing Values

We will now discuss a few simple approaches to deal with missing values

- We cannot easily assess them on our traffic data
- ...Since we do not know the ground truth of each missing value

Therefore, we will initially use partially synthetic data

- We will start from a time series without any missing value, then remove values artificially
- ...And measure the accuracy of our filling approaches via the Root MSE


$$RMSE = \sqrt{\frac{1}{n} \sum_{i=0}^n (\tilde{x}_i - \hat{x}_i)^2}$$

Where \tilde{x}_i is a value from the filled series and \hat{x}_i the ground truth

The Benchmark Dataset

Our benchmark dataset consists of twitter volume related to amazon.com

```
In [2]: nab.plot_series(data1, labels1, windows1, figsize=figsize)
```



- There are a few anomalies, but we are not concerned with them
- Then series has a seasonal/periodic component

The Benchmark Dataset

We now introduce some missing values

First, we draw some starting points at random

```
In [3]: np.random.seed(42) # seed (to get reproducible results)
mv_num = 30 # number of intervals with missing values
mv_starts = np.random.choice(range(len(data1.index)), mv_num, replace=False)
mv_starts.sort()
mv_starts[:4]
```

```
Out[3]: array([ 88, 123, 169, 304])
```

Then, we clear values over increasingly long intervals:


```
In [4]: data1_mv = data1.copy()
for i in range(mv_num):
    data1_mv.iloc[mv_starts[i]:mv_starts[i]+i+1] = np.NaN
```

- The first interval contains 1 missing value, the second 2, the third 3...

The Benchmark Dataset

Let's have a look at the results around one of the "holes"

```
In [7]: idx, pad = 10, 20  
nab.plot_series(data1_mv.iloc[mv_starts[idx]-pad:mv_starts[idx]+mv_num+pad+1], figsize=figsize)  
plt.plot(data1_mv.iloc[mv_starts[idx]-1:mv_starts[idx]+idx+2]);
```



- The orange part corresponds to the removed values

Forward/Backward Filling

The easiest approach for missing values consists in **replicating nearby observations**

- **Forward filling:** propagate forward the last valid observation
- **Backward filling:** propagate backward the next valid observation

An important observation:

- When filling missing values, **we have access to the whole series**
- ...So we can reason **both forward and backwards**

Forward/backward filling are simple methods, but they can work well

- Rationale: most time series have a certain "inertia"
- ...I.e.: a strong level of local correlation
- For this reason, the last observation is a good predictor for the next one
- ...Remember the **persistence** predictor?

Forward/Backward Filling

Forward and backward filling are pre-implemented in pandas

They are available through the `fillna` method:

```
DataFrame.fillna(..., method=None, ...)
```

- `fillna` replaces NaN values in a DataFrame Or Series
- The `method` parameter can take the values:
 - "pad" or "ffill": these correspond to forward filling
 - "backfill" or "bfill": these correspond to backward filling

They are generally applied to datasets with a dense index

- Remember that our benchmark dataset already has a dense index

Forward/Backward Filling on the Benchmark

We can finally test forward/backward filling

```
In [8]: ffseries = data1_mv.fillna(method='ffill')
        bfseries = data1_mv.fillna(method='bfill')
```

We can check the corresponding MSE:

```
In [9]: from sklearn.metrics import mean_squared_error
        rmse_ff = np.sqrt(mean_squared_error(data1, ffseries))
        rmse_bf = np.sqrt(mean_squared_error(data1, bfseries))
        print(f'RMSE for forward filling: {rmse_ff:.2f}, for backward filling {rmse_bf:.2f}')
```


```
RMSE for forward filling: 3.18, for backward filling 6.45
```

- In this case forward filling seems to work better
- The results are of course application-dependent

Forward/Backward Filling on the Benchmark

Let's have a close look at forward filling around some of the missing values

```
In [10]: idx, pad = 0, 20  
nab.plot_series(data1.iloc[mv_starts[idx]-pad:mv_starts[idx]+mv_num+pad+1], figsize=figsize)  
plt.plot(ffseries.iloc[mv_starts[idx]-1:mv_starts[idx]+idx+2]);
```




- This is the first (and shortest) gap

Forward/Backward Filling on the Benchmark

Let's have a close look at forward filling around some of the missing values

```
In [11]: idx, pad = mv_num//2, 20  
nab.plot_series(data1.iloc[mv_starts[idx]-pad:mv_starts[idx]+mv_num+pad+1], figsize=figsize)  
plt.plot(ffseries.iloc[mv_starts[idx]-1:mv_starts[idx]+idx+2]);
```




- This is an intermediate-length gap

Forward/Backward Filling on the Benchmark

Let's have a close look at forward filling around some of the missing values

```
In [12]: idx, pad = mv_num-1, 20  
nab.plot_series(data1.iloc[mv_starts[idx]-pad:mv_starts[idx]+mv_num+pad+1], figsize=figsize)  
plt.plot(ffseries.iloc[mv_starts[idx]-1:mv_starts[idx]+idx+2]);
```



■ This is the longest gap

Forward/Backward Filling on the Traffic Data

Let's now consider the traffic data

We will limit ourselves to forward filling:

```
In [13]: ddata2 = data2.resample('5min').mean()
         ddata2_ff = ddata2.fillna(method='ffill')
```

- data2 contains the traffic series

We now isolate the filled values

This is needed so that we can highlight them in the forthcoming plots


```
In [14]: ddata2_ffonly = ddata2_ff.copy()
         ddata2_ffonly[~ddata2['value'].isnull()] = np.nan
```

- ffseries and bfseries now contain **only** the filled values

Forward/Backward Filling on the Traffic Data

Let us have a look at the results of forward filling:


```
In [16]: nab.plot_series(ddata2_ff, show_markers=True, figsize=figsize)
plt.scatter(ddata2_ffonly.index, ddata2_ffonly, marker='.', c='tab:orange');
```



Forward/Backward Filling on the Traffic Data

Forward filling works reasonably for small gaps


```
In [17]: mask = (ddata2_ff.index >= '2015-09-02 01:00') & (ddata2_ff.index < '2015-09-03 00:00')
nab.plot_series(ddata2_ff[mask], show_markers=True, figsize=figsize)
plt.scatter(ddata2_ffonly.index[mask], ddata2_ffonly[mask], marker='.', c='tab:orange');
```



Forward/Backward Filling with our Series

...But it is not particularly effective for larger gaps

```
In [18]: mask = (ddata2_ff.index >= '2015-09-08 18:00') & (ddata2_ff.index < '2015-09-09 18:00')
nab.plot_series(ddata2_ff[mask], show_markers=True, figsize=figsize)
plt.scatter(ddata2_ffonly.index[mask], ddata2_ffonly[mask], marker='.', c='tab:orange');
```



(Geometric) Interpolation

A few more options are available via the interpolate method

```
DataFrame/Series.interpolate(method='linear', ...)
```

The `method` parameter determines how NaNs are filled:

- "linear" uses a linear interpolation, assuming uniformly spaced samples
- "time" uses a linear interpolation, but supports non-uniformly spaced samples
- "nearest" uses the closest value
- "polynomial" uses a polynomial interpolation
- Even "ffill" and "bfill" are available

Both "polynomial" and "spline" require to specify the additional parameter `order`

- E.g. `df.interpolate(method='polynomial', order='3')`

(Geometric) Interpolation

Let us check the performance of some approaches

```
In [19]: args = [{'method': 'linear'}, {'method': 'time'}, {'method': 'nearest'},  
                {'method': 'polynomial', 'order': 2}]  
  
for a in args:  
    tmp = data1_mv.interpolate(**a)  
    rmse = np.sqrt(mean_squared_error(data1, tmp))  
    print(f'RMSE for {a["method"]}: {rmse:.2f}')
```

```
RMSE for linear: 4.42  
RMSE for time: 4.42  
RMSE for nearest: 4.95  
RMSE for polynomial: 17.89
```


- "linear" and "time" are equivalent in this case (we have uniformly-spaced samples)
 - "polynomial" is the most complex, and in this case also the worst
- All perform **worse** than forward filling (at least in this case)!

Curve Interpolation on Traffic Data

On real data, we cannot (easily) compare interpolation methods

- This is because typically we have **no ground truth**
- For now, we will settle for a visual inspection

```
In [20]: ddata2_ln = ddata2.interpolate(method='linear')  
nab.plot_series(ddata2, filled_version=ddata2_ln, figsize=figsize)
```



Considerations

All these methods for dealing with missing values:

- Work ok for small gaps, but loose effectiveness as the gap size grows
- This is true **even for the more advanced** filling methods!
- Better methods just degrade a bit more slowly

MSE comparisons can be very effective:

- But requires access to ground truth!
- An idea: **make** your ground truth by artificially removing values
- ...Then comparing filling methods based on their performance on the artificial gaps

MSE is not everything

- Every filling method makes mistakes: it's important not to make the **wrong ones**
- Think about how you plan to use your (filled) time series
- Can you expect the series to have uniform variance?