

Gaussian Processes

Filling Values Using a Model

If we want to fill missing values better, we need to **predict them**

We need a **model**, which can infer their value

- All the approaches seen so far can be considered (extremely simple) models
- ...We just need a more advanced one!

What are the **desired properties of the model we seek?**

Given a gap (i.e. one or more contiguous missing values), the model:

- Must be able to make a prediction about the missing values
- ...Which is consistent with all the available observations
- I.e. it should be able to **interpolate** the data (in generalized sense)

Most ML models **cannot be used for filling (in a straightforward fashion)**

Filling Values Using a Model

Density estimation does not (natively) provide predictions

To be fair, predictions can be **extracted** from a density estimator:

- Given an estimator $f(\mathbf{x}, \theta)$ for $P(\mathbf{x})$ we can find the most likely value for \mathbf{x} by solving:

$$\operatorname{argmax}_{\mathbf{x}} f(\mathbf{x}, \theta)$$

- This is a **Maximum A Posteriori (MAP)**
- ...And its what most regressors/classifiers natively compute

However, with a density estimator, computing the MAP can be **very expensive**


Auto-regressors makes use of past observations, but not the future ones

- They are designed for **extrapolation** (predict beyond the boundaries)
- ...And not for **interpolation**

Gaussian Processes

One of the few viable ML models is given by **Gaussian Processes (GP)**

We will introduce their key concepts via an example:




- Say we want to model how rainfall changes over a stretch of land
- $y = \text{rainfall}$, $(x_0, x_1) = \text{position on the surface of land}$

Gaussian Processes

One of the few viable ML models is given by **Gaussian Processes (GP)**

Since this is a physical phenomenon...




- ...We can reasonably assume that \mathbf{y} is Normally distributed
- But unless we know more, we can say nothing else

Gaussian Processes

One of the few viable ML models is given by **Gaussian Processes (GP)**

However, if we have a few measurements...




We can assume that **rainfall in nearby locations is similar**

- We can rely on this to estimate **the chance of the measurements themselves**
- ...And of rainfall at **positions for which we lack measurements**

Gaussian Processes

One of the few viable ML models is given by **Gaussian Processes (GP)**

We view the measurements as components of a variable $\mathbf{y}_X = (y_0, y_1, y_2)$



- \mathbf{y}_X will follow a multivariate Normal distribution
- ...And the covariance will depend on the distance between measurements

Gaussian Processes

Formally:

A GP is a **stochastic process**, i.e. a collection of indexed random variables

- Each variable y_x is indexed via a tuple x (e.g. location, time...)
- The index is **continuous** and the collection **infinite**
- Every finite subset of y_x variables follows a **Multivariate Normal Distribution**

Some examples:

- y_x could be the rainfall rate at location x
- y_x could be the twitter volume at time x
- y_x could be the traffic volume at time x

In general y_x is the value of a (stochastic) function for input x

Multivariate Normal Distribution

The multivariate normal distribution?

- It works for many real world phenomena
- It has a closed-form density function that is (relatively) easy to compute
- The PDF is defined via a (vector) mean μ and a covariance matrix Σ
- ...And often we can assume $\mu = \mathbf{0}$, so knowing Σ is enough

So, given a set of indexes/input values of interest X

Then, if we manage to know Σ , we can compute:

- The probability density $f(\hat{\mathbf{y}}_X)$ of some given observations
 - I.e. the probability of a dataset
- The conditional density $f(\hat{\mathbf{y}}_x \mid \hat{\mathbf{y}}_X)$ of a new observation
 - I.e. a prediction w.r.t. a set of known observations $\hat{\mathbf{y}}_X$

Defining the Covariance Matrix

How do we define Σ ?

We assume that the covariance depends on \mathbf{x} (and not on the \mathbf{y})

- Given two variables \mathbf{y}_{x_i} and \mathbf{y}_{x_j} , their covariance is given by $K(x_i, x_j)$
- Where K is called a kernel function (and is user chosen)

Given any finite set of variables $\{\mathbf{y}_{x_1}, \dots, \mathbf{y}_{x_n}\}$, the covariance matrix is:

$$\Sigma = \begin{pmatrix} K(x_1, x_1) & K(x_1, x_2) & \dots & K(x_1, x_n) \\ K(x_2, x_1) & K(x_2, x_2) & \dots & K(x_2, x_n) \\ \vdots & \vdots & \vdots & \vdots \\ K(x_n, x_1) & K(x_n, x_2) & \dots & K(x_n, x_n) \end{pmatrix}$$

- I.e. it's entirely specified via the kernel

Unfortunately, choosing the kernel completely by hand would still be too difficult

Fitting a Gaussian Process

In practice, to define the kernel we:

- Pick a **parameterized** kernel function $K_{\theta}(x_i, x_j)$, where θ = parameter vector
- Collect training observations $\hat{\mathbf{y}}_{\mathbf{X}}$

Then we choose θ so as to maximize the **likelihood** of the training data, i.e.:

$$\operatorname{argmax}_{\theta} f(\hat{\mathbf{y}}_{\mathbf{X}}, \theta)$$

- Where $f(\hat{\mathbf{y}}_{\mathbf{X}}, \theta)$ is the **estimated** probability density of the observations

The training problem


- Is a (possibly challenging) numerical optimization problem
- ...Which is typically solved to **local optimality** (e.g. via gradient descent)

Gaussian Processes in scikit-learn

Let's see how to use Gaussian Processes in scikit-learn

First, let us choose a target function:

```
In [2]: f = lambda x: x * np.sin(2*np.pi*x) + x # target function  
x = np.linspace(0, 3, 1000)  
y = pd.Series(index=x, data=f(x))  
nab.plot_gp(target=y, figsize=figsize)
```




Gaussian Processes in scikit-learn

Let's see how to use Gaussian Processes in scikit-learn

Then we build a small training set:

```
In [3]: np.random.seed(42)
n_tr = 15
x_tr = np.linspace(0.2, 2.8, n_tr) + 0.2*np.random.rand(n_tr)
x_tr.sort()
y_tr = pd.Series(index=x_tr, data=f(x_tr) + 0.2*np.random.rand(n_tr))
nab.plot_gp(target=y, samples=y_tr, figsize=figsize)
```



Gaussian Processes in scikit-learn

Let's see how to use Gaussian Processes in scikit-learn

Then we need to choose a kernel

- There are many available options
- We will **start** with a simple Radial Basis Function (i.e. Gaussian) kernel

$$K(x_i, x_j) = e^{-\frac{d(x_i, x_j)^2}{2l}}$$

The correlation **decreases with the (Euclidean) distance** $d(x_i, x_j)$:

- Intuitively, **the closer the points, the higher the correlation**
- The l parameter (**scale**) control the rate of the reduction

We have 15 indexes, so the kernel will define a 15×15 covariance matrix

Gaussian Processes in scikit-learn

Here's how to use an RBF kernel in scikit-learn

```
In [4]: from sklearn.gaussian_process.kernels import RBF  
  
kernel = RBF(1, (1e-3, 1e3))
```

The RBF kernel has a single parameter, representing its **scale**

The extra (tuple) parameter represents a pair of **bounds**

- During training only parameter values within the boundaries will be considered
- Bounds can be very useful for controlling the training process
- ...Based on the available domain information

Gaussian Processes in scikit-learn

Now we can train a Gaussian Process

```
In [5]: from sklearn.gaussian_process import GaussianProcessRegressor
gp = GaussianProcessRegressor(kernel=kernel, n_restarts_optimizer=9)
gp.fit(y_tr.index.values.reshape(-1,1), y_tr.values) # needs 2D input
gp.kernel_
```

```
Out[5]: RBF(length_scale=0.229)
```

- Training uses Gradient Descent...
- ...To maximize the likelihood (density) of the training data
- **Restarts** are needed to mitigate issues due to local optima


And then we can obtain the predictions:

```
In [6]: xp, std = gp.predict(x.reshape(-1,1), return_std=True)
xp = pd.Series(index=y.index, data=xp)
std = pd.Series(index=y.index, data=std)
```


Gaussian Processes in scikit-learn

We can now plot the predictions

```
In [7]: nab.plot_gp(target=y, samples=y_tr, pred=yp, std=std, figsize=figsize)
```



- We get both a **point estimate** (the blue line)
- ...And a **confidence interval** (the light blue areas)

But how did we manage that?

Behind the Scenes

There are **two tricks at work** here:

1) The training examples $\hat{\mathbf{y}}_{\mathbf{X}}$ are **part of the GP parameters**

- This is similar to what we have (e.g.) in KDE
- ...And quite different from other ML methods (e.g. NN, DTs, SVMs...)

Given an new input value \mathbf{x} , the GP output corresponds to:

$$f(y_x \mid \hat{\mathbf{y}}_{\mathbf{X}})$$

- I.e. the probability of y_x **conditioned on the known examples**

2) For computing Σ , we only need to know the **input values** (i.e. \mathbf{X} and \mathbf{x})

- This is by construction: all our kernels $K(\mathbf{x}_i, \mathbf{x}_j)$ are built this way
- Which means that we can compute $f(y_x \mid \hat{\mathbf{y}}_{\mathbf{X}})$ via a closed-form expression

Behind the Scenes

Thanks to this, for $f(y_x \mid \hat{\mathbf{y}}_{\mathbf{X}})$ we can get

The **conditional mean** (which is also the MAP):

$$\arg \max_{y_x} f(y_x \mid \hat{\mathbf{y}}_{\mathbf{X}})$$

- This will be our "prediction"

The **conditional standard deviation**

$$\sqrt{\text{Var} [f(y_x \mid \hat{\mathbf{y}}_{\mathbf{X}})]}$$

- This is used to define the confidence intervals

In practice, we GP output is a probability distribution

A Numeric Example

As an example, say we want a prediction for $x = 2.5$, i.e. $y_{2.5}$

...And let's assume that our training set contains **only one example**

- We will consider **separately** the tenth and first example in our dataset
- We have: $(\hat{x}_9, \hat{y}_{\hat{x}_9}) \simeq (2.01, 2.27)$ and $(\hat{x}_0, \hat{y}_{\hat{x}_0}) \simeq (0.27, 0.58)$

The covariance matrix in the two cases is therefore:

$$\Sigma_{y_x, \hat{y}_{\hat{x}_9}} = \begin{pmatrix} K(2.01, 2.01) & K(2.01, 2.5) \\ K(2.5, 2.01) & K(2.5, 2.5) \end{pmatrix}$$

$$\Sigma_{y_x, \hat{y}_{\hat{x}_0}} = \begin{pmatrix} K(0.27, 0.27) & K(0.27, 2.5) \\ K(2.5, 0.27) & K(2.5, 2.5) \end{pmatrix}$$

- Each matrix defines a multivariate Normal distribution

Behind the Scenes

Let's actually build the matrices in Python

- Note: scikit-learn kernels are not designed to be used on individual points
- So, for this we will rely on basic numpy methods

We start with \hat{x}_9 and x , which are **close to each other**

```
In [8]: from scipy.stats import multivariate_normal
X9, X0, X = [[x_tr[9]]], [[x_tr[0]]], [[2.5]] # Must be 2D
sigma_9x = np.array([[kernel(X9, X9)[0,0], kernel(X9, X)[0,0]],
                    [kernel(X, X9)[0,0], kernel(X, X)[0,0]]])
f_9x = multivariate_normal([0, 0], cov=sigma_9x)
```


Then we do the same for \hat{x}_0 and x , which are **far apart**

```
In [9]: sigma_0x = np.array([[kernel(X0, X0)[0,0], kernel(X0, X)[0,0]],
                             [kernel(X, X0)[0,0], kernel(X, X)[0,0]]])
f_0x = multivariate_normal([0, 0], cov=sigma_0x)
```

Behind the Scenes

\hat{x}_9 and x are **close to each other**, so $\hat{y}_{\hat{x}_9}$ and y_x are **strongly correlated**

```
In [10]: yr = np.linspace(-5, 5, 100)
nab.plot_distribution_2D(f_9x, yr, yr, figsize=figsize)
plt.xlabel('y_{x_9}'); plt.ylabel('y_{x}'); plt.tight_layout()
```




- Still, if we know **neither** $\hat{y}_{\hat{x}_9}$ nor y_x , we can only say that they are likely **both zero**

Behind the Scenes

But we **do know** $\hat{y}_{\hat{y}_9}$! So, we can use this information_

```
In [11]: nab.plot_distribution_2D(f_9x, yr, yr, figsize=figsize)
plt.axvline(10*(y_tr[x_tr[9]] + 5), color='tab:orange');
plt.xlabel('y_{x_9}'); plt.ylabel('y_{x}'); plt.tight_layout()
```




- Given the observation, the most likely value for y_x changes considerably

Behind the Scenes

\hat{x}_0 and x are **far apart**, so $\hat{y}_{\hat{x}_0}$ and y_x are **loosely correlated**

```
In [12]: nab.plot_distribution_2D(f_0x, yr, yr, figsize=figsize)
plt.axvline(10*(y_tr[x_tr[0]] + 5), color='tab:orange');
plt.xlabel('y_{x_0}'); plt.ylabel('y_{x}'); plt.tight_layout()
```



- Knowing $\hat{y}_{\hat{x}_0}$ is not going to be of much help here

Behind the Scenes

So, a few key insight to keep in mind:

- Superficially, GPs behave like functions that output probability distribution
- Internally, this is enabled by two components:
 - The kernel, defining how all the points are correlated
 - A set of observations, use to obtain conditional distributions

In scikit-learn:

When we call the `fit` method:

- The optimizer adjusts the kernel parameters
- ...And the observations \hat{y}_x are stored

When we call the `predict` method:


- The covariance matrix is built
- The model computes the conditional distributions

How to Improve the Model

Now, let's try to improve our model

We need to choose a kernel appropriate to our dataset

```
In [13]: nab.plot_gp(target=y, samples=y_tr, figsize=figsize)
```



- We have a bit of noise, a period, and a trend

How to Improve the Model

So, let us deal with the noise first

```
In [14]: from sklearn.gaussian_process.kernels import WhiteKernel  
  
kernel = WhiteKernel(0.1, (1e-3, 1e3))  
kernel += RBF(1, (1e-2, 1e2))
```

WhiteKernel captures the presence of noise in the data

$$K(x_i, x_j) = \sigma^2 \text{ iff } x_i = x_j, 0 \text{ otherwise}$$

- The only parameter of `WhiteKernel` represents the noise level σ^2
- A small noise level prevent overfitting the data
- ...But too much noise leads to useless predictions!

How to Improve the Model

It's often a good idea to have magnitude parameters in the kernel

```
In [15]: from sklearn.gaussian_process.kernels import ConstantKernel

kernel = WhiteKernel(0.1, (1e-2, 1e2))
kernel += ConstantKernel(1, (1e-2, 1e2)) * RBF(1, (1e-2, 1e2))
```

ConstantKernel is a constant factor (in this case a relative weight)

■ ...And allows the optimizer to tune the magnitude of the RBF kernel

Let's repeat training again:

```
In [16]: gp = GaussianProcessRegressor(kernel=kernel, n_restarts_optimizer=9)
gp.fit(y_tr.index.values.reshape(-1,1), y_tr.values) # needs 2D input
print(gp.kernel_)
```


```
WhiteKernel(noise_level=0.01) + 2.21**2 * RBF(length_scale=0.321)
```

```
/usr/local/lib/python3.9/site-packages/sklearn/gaussian_process/kernels.py:402: ConvergenceWarning: The optimal value found for dimension 0 of parameter k1__noise_level is close to the specified lower bound 0.01. Decreasing the bound and calling fit again may find a better value.
  warnings.warn("The optimal value found for "
```

How to Improve the Model

Let us see the new predictions

```
In [17]: xp, std = gp.predict(x.reshape(-1,1), return_std=True)
xp = pd.Series(index=y.index, data=xp)
std = pd.Series(index=y.index, data=std)
nab.plot_gp(target=y, samples=y_tr, pred=xp, std=std, figsize=figsize)
```



- Better, but we are still not exploiting the period and the trend

How to Improve the Model

So, let us take them into account, starting with the period

```
In [18]: from sklearn.gaussian_process.kernels import ExpSineSquared
kernel = WhiteKernel(0.1, (1e-2, 1e2))
kernel += ConstantKernel(1, (1e-2, 1e2)) * RBF(1, (1e-2, 1e2))
kernel += ExpSineSquared(1, 1, (1e-2, 1e2), (1e-2, 1e2))
```

ExpSineSquared captures the period:

$$K(x_i, x_j) = e^{-2 \frac{\sin^2\left(\pi \frac{d(x_i, x_j)}{p}\right)}{l^2}}$$

- The correlation grows as the distance is close to a multiple of the period p
- The scale parameter l controls the rate of decrease/increase
- In the implementation, the first parameter is l and the second p

How to Improve the Model

Now, let's try to capture the trend

```
In [19]: from sklearn.gaussian_process.kernels import DotProduct
kernel = WhiteKernel(0.1, (1e-2, 1e2))
kernel += ConstantKernel(1, (1e-2, 1e2)) * RBF(1, (1e-2, 1e2))
kernel += ExpSineSquared(1, 1, (1e-2, 1e2), (1e-2, 1e2))
kernel += DotProduct(1, (1e-2, 1e2))
```

DotProduct (somewhat) captures the trend:

$$K(x_i, x_j) = \sigma^2 + x_i x_j$$

- The larger the \mathbf{x} values, the larger the correlation
- This allows the distance from the mean (which is zero) to grow
- The σ parameter controls the base level of correlation
- Unlike all kernels so far DotProduct is **not translation-invariant** (we sa!

How to Improve the Model


The new predictions are a bit better at the edges of the plot

```
In [20]: gp = GaussianProcessRegressor(kernel=kernel, n_restarts_optimizer=9)
gp.fit(y_tr.index.values.reshape(-1,1), y_tr.values) # needs 2D input
print(gp.kernel_)
xp, std = gp.predict(x.reshape(-1,1), return_std=True)
xp = pd.Series(index=y.index, data=xp)
std = pd.Series(index=y.index, data=std)
nab.plot_gp(target=y, samples=y_tr, pred=xp, std=std, figsize=figsize)
```

```
/usr/local/lib/python3.9/site-packages/sklearn/gaussian_process/kernels.py:402: ConvergenceWarning: The optimal value found for dimension 0 of parameter k1__k1__k1__noise_level is close to the specified lower bound 0.01. Decreasing the bound and calling fit again may find a better value.
```

```
warnings.warn("The optimal value found for "
```

```
WhiteKernel(noise_level=0.01) + 1.17**2 * RBF(length_scale=0.305) + ExpSineSquared(length_scale=2.17, periodicity=0.939) + DotProduct(sigma_0=0.0197)
```



Considerations

Gaussian Processes are a **very flexible** ML technique

- They can be used to making predictions
- ...Together with their confidence intervals
- ...But also for (conditional) density estimation
- ...And for generating data

They are **non-trivial to use**:

- In particular, choosing a kernel requires some practice and some understanding
- Automating the process is possible, but complex (grid search is likely not enough)

Gaussian processes tend to perform better: